



Approximate Bayesian Computation for infectious disease modelling

Amanda Minter^{a,1,*}, Renata Retkute^{b,1}

^a Centre for the Mathematical Modelling of Infectious Diseases, London School of Hygiene and Tropical Medicine, London, UK

^b The Zeeman Institute for Systems Biology and Infectious Disease Epidemiology Research, University of Warwick, UK

ARTICLE INFO

Keywords:

Approximate Bayesian Computation

R

Epidemic model

Stochastic model

Spatial model

ABSTRACT

Approximate Bayesian Computation (ABC) techniques are a suite of model fitting methods which can be implemented without a using likelihood function. In order to use ABC in a time-efficient manner users must make several design decisions including how to code the ABC algorithm and the type of ABC algorithm to use. Furthermore, ABC relies on a number of user defined choices which can greatly effect the accuracy of estimation. Having a clear understanding of these factors in reducing computation time and improving accuracy allows users to make more informed decisions when planning analyses. In this paper, we present an introduction to ABC with a focus of application to infectious disease models. We present a tutorial on coding practice for ABC in R and three case studies to illustrate the application of ABC to infectious disease models.

1. Introduction

Mathematical models can be used to predict infectious disease dynamics at multiple scales. Model fitting, the process of estimating the parameters of the mathematical model from data, can be performed using a number of different methods. For infectious disease models, researchers are often confronted with missing data due to the nature of partially observed epidemics. Methods are needed to fit these models to data in an accurate and time efficient manner. In the Bayesian framework model parameters are assumed to be random variables and therefore have probability distributions, we seek to estimate the posterior distribution of these parameters (Gelman et al., 2013).

Approximate Bayesian Computation (ABC) methods can be used to approximate these posterior distributions when a likelihood function is intractable or not known (Beaumont et al., 2002), and can be used when the data available are coarse or complex. The ABC rejection algorithm Perez-Lezaun et al. (1999) is the most straight-forward of ABC methods yet ABC methods often yield inefficient sampling of the parameter space (Sadegh and Vrugt, 2014). As such, other algorithms with more efficient sampling schemes have been developed. The ABC-Sequential Monte Carlo (ABC-SMC) algorithm, provides a computationally efficient estimation procedure compared to traditional ABC rejection algorithms (Toni et al., 2009).

ABC is a powerful tool for model fitting and has a lot of flexibility with its application. However, this also means that the user must be

able to defend, and is responsible for all of these choices. There is an existing body of literature on ABC tutorials. In particular, for an starting introduction to ABC in general Sunnåker et al. (2013), for further ABC examples, including model selection see Toni et al. (2009), for tutorial on ABC for stochastic epidemic models Kyraios et al. (2017), Hartig et al. (2002), McKinley et al. (2018), for comparison of likelihood-based Markov Chain Monte Carlo and ABC in epidemic modelling McKinley et al. (2009), for examples for population genetics see Marjoram et al. (2003), Beaumont et al. (2002).

The purpose of this paper is to provide a tutorial on ABC for infectious disease models for readers with very little experience in model fitting. To that end, we provide a gentle introduction to the concept from the Bayesian inference paradigm followed by a presentation of how the ABC algorithm can be implemented in R. We present the implementation of ABC using three case studies to highlight the different approaches that can be taken to perform model fitting using ABC. Case study 1 is an introductory example which illustrates the application of the ABC-rejection algorithm to a deterministic epidemic model, with a particular focus on the power of ABC based on the model to be fitted and the data available.

Case study 2 is a stochastic compartmental model with age-structure, i.e. the model where population is divided into compartments classified by the infection status and age of the host (Anderson and May, 1991). The dynamics of many childhood infections such as measles (Babad et al., 1995), whooping cough (Campbell et al., 2015),

Abbreviations: ABC, Approximate Bayesian Computation; SMC, Sequential Monte Carlo; NN, nearest neighbour

* Corresponding author.

E-mail address: amanda.minter@lshtm.ac.uk (A. Minter).

¹ These authors contributed equally to this work..

<https://doi.org/10.1016/j.epidem.2019.100368>

Received 22 February 2019; Received in revised form 20 August 2019; Accepted 30 August 2019

Available online 25 September 2019

1755-4365/ © 2019 Published by Elsevier B.V. This is an open access article under the CC BY-NC-ND license

(<http://creativecommons.org/licenses/by-nc-nd/4.0/>).

chickenpox (Schuette, 1999) and rubella (Kanaan and Farrington, 2005) can be described by age-structured compartmental model. In this case study we will investigate the 2010 measles outbreak in Malawi (Minetti et al., 2013) with the data consisting of weekly cases and percentage of cases in three age groups. Case study 3 uses an individual-based model which accounts for the spatial interaction between individual hosts (Keeling and Rohani, 2007). These models have been used to investigating the spatial dynamics of foot-and-mouth disease (Keeling et al., 2001), avian influenza (Hill et al., 2017) or transmission of visceral leishmaniasis (Chapman et al., 2018). In this case study we will investigate citrus tristeza virus spread in an orchard (Marcus et al., 1984) with data consisting of maps of disease incidence observed at discrete times.

2. Approximate Bayesian Computation

The paradigm of Bayesian inference is based on the idea of updating belief with new evidence (Gelman et al., 2013). If we denote θ our mathematical model parameters and D our data, then we can use Bayes rule to determine the posterior distribution of the parameters given the data, $P(\theta|D)$:

$$P(\theta|D) = \frac{P(D|\theta)P(\theta)}{P(D)} \quad (1)$$

$$\propto P(D|\theta)P(\theta) \quad (2)$$

where $P(\theta)$ is the prior distribution (which represents our prior belief) and $P(D|\theta)$ is the likelihood function, the probability density function for the data given the parameters. The second equation is written without the marginal likelihood of the data, $P(D)$, as the value of $P(D)$ does not depend on θ . Hence we can say that $P(\theta|D)$ is proportional to the numerator of Eq. (1).

ABC is a method for approximating the posterior probability $P(\theta|D)$ without a using likelihood function (Beaumont, 2010). To illustrate this, we will first present a basic ABC algorithm, the ABC rejection algorithm (Perez-Lezaun et al., 1999; Toni et al., 2009). The algorithm ‘accepts’ proposed parameter values, so called particles, based on whether the distance $d()$ between the data D and the model simulated data D^* is less than or equal to some threshold ϵ . The steps of the ABC-rejection algorithm are as follows:

1. Sample θ^* from the prior distribution $P(\theta)$.
2. Simulate data set from the model, using parameters θ^* , to get D^* .
3. If $d(D, D^*) \leq \epsilon$ accept θ^* otherwise reject.
4. Repeat until N particles (the parameter values or parameter sets) $\Theta^* = \{\theta_i^*; i = 1, \dots, n\}$ are accepted.

where $d(D, D^*)$ is some distance measure between the model output and the data, and ϵ is a tolerance value which is chosen by the user. The N accepted samples provide an approximation of the posterior distribution of the model parameters θ ,

$$P(\theta|D) \approx P(\theta|d(D, D^*) \leq \epsilon) \quad (3).$$

The distance measure $d(D, D^*)$ implies it is possible to directly compare the model output to the data. For example, if the data are daily cases of an infection and we wish to estimate parameters of a Susceptible-Infected-Recovered model (SIR), the user could compare the sum of the squared differences between the number of infected individuals between the data and that which is predicted by the model over time i.e. $d(D, D^*) = \sum_{t=1}^T (D_t - D_t^*)^2$ where time is denoted by $t = 0, \dots, T$.

When working with infection dynamics, there is often missing data or the data may consist of one measure of an epidemic. For example, say the data consists of the final size of the epidemic which we denote μ . Then we simulate from the underlying model to calculate the final size μ^* from our simulated data set D^* using a ‘summary statistic’ $S(D^*)$. In

this example the ‘summary’ is the calculation of the final size of the epidemic but it can take a form of other aggregated measures such as the number of infected individuals in different age groups, or the number of infected individuals in particular geographical locations. Then step 3 in the ABC-rejection algorithm becomes,

1. If $d(\mu, \mu^*) \leq \epsilon$, where $\mu^* = S(D^*)$, accept θ otherwise reject.

An intersection metric can be used when there is more than one summary statistic (McKinley et al., 2018):

$$\hat{P}(D|D^*) = \prod_m \mathbb{1}(d^m(D, D^*) \leq \epsilon^m), \quad (4)$$

where $\mathbb{1}$ is a binary operator and m is a number of summary statistics. The intersection metrics ensures that model run is accepted only if all criteria are satisfied, but requires the specification of multiple tolerance values at each generation.

2.1. Implementing ABC in R

In this section we illustrate how to implement an efficient ABC-rejection algorithm in R. To implement ABC in R, there are a number of techniques that can be used for fast computation. We direct the readers to Visser et al. (2015), Wilson et al. (2014) for tutorials on efficient coding. For ABC-rejection algorithms, computation time can be decreased by pre-allocation of memory to store the accepted particles, memorization (do not repeat calculations with the loop) of data calculations, not printing unnecessarily, and vectorisation where possible in the function calls.

Fig. 1 shows the code for the ABC-rejection algorithm implemented in R. There are two distinct sections to the code: the initial set up which occurs outside of the while loop (the initialisation), and the ABC-rejection algorithm which occurs in the while loop. The initialisation requires the user to specify inputs related to the ABC-rejection algorithm. Specifically, the desired number of particles (N) and the tolerance value (ϵ). The user then needs to pre-allocate memory to store the results of the ABC-rejection algorithm. For a model with one parameter this would be a vector, if the model had multiple parameters then this would be a matrix. In addition, the user may wish to store the distances which were accepted, which would require an additional column in the matrix.

The ABC-rejection algorithm is performed in a while loop with a counter (i) initiated outside of the loop. Within the while loop, we utilise existing R functions to propose values from the prior distributions (for example `runif`, `rnorm`). Then we recommend user coded functions to separate the ABC-rejection steps. Data are simulated using the epidemic model and with the proposed parameters (θ^*) using the user coded function `run_model`. The distance between the data and the model output is calculated using the user coded function `calc_distance`. If the distance is less than or equal to the tolerance, `epsilon`, the current parameter value is stored and the counter is updated. Additionally, if a summary statistic is to be used the user can code a function to calculate the statistic using the simulated data set. The loop continues until the desired number of accepted particles (N) have been found. The user could specify a tolerance value that is too small for any proposed value to be accepted which could result in an infinite while loops. To avoid this the user could define a maximum number of iterations to be executed.

The most computationally efficient ABC-rejection algorithm implemented in R will still be slow. The algorithm proposes parameter values from the prior at each iteration, and so does not remain in parameter space with small distances between the data and model simulation. However, due to the lack of dependencies between iterations, the ABC rejection algorithm can be easily parallelised Baragatti and Pudlo (2014), wherein the number of desired particles can be divided up, ran independently and then combined into one large set of particles,

```

#### ABC set up ####

N <- 1000 # Number of accepted particles
epsilon <- 20 # Epsilon value
n_par <- 1 # How many parameters will be estimated
res <- matrix(ncol = n_par + 1, nrow = N) # Empty matrix to store results
# model parameter(s) and the distance

#### ABC algorithm ####

i <- 1 # Initiate counter of accepted particles
j <- 1 # Initiate counter of proposed particles
while(i <= N){ # While the number of accepted particles is less than N

  theta_star <- runif(1, 0, 1) # Sample from prior distribution(s)
  D_star <- run_model(theta_star) # Simulate data set from the model
  distance <- calc_distance(D, D_star) # Calculate distance

  if(distance <= epsilon){ # If the distance is less than or equal to the tolerance
    res[i, ] <- c(theta_star, distance) # Store results
    i <- i + 1 # Update counter
  }

  j <- j + 1 # Update counter
  acc_rate <- i / j # Calculate the acceptance rate
  cat("current acceptance rate = ", round(acc_rate, 2), "\n")
}

```

Fig. 1. Code for an of ABC-rejection algorithm implemented in R.

so-called embarrassingly parallel. Ultimately, the ABC-rejection algorithm is slow due to its inefficient proposal of parameter values. In the next section, we introduce an extension to the ABC-rejection algorithm which uses a more efficient scheme to sample the parameter space.

2.2. Speeding up ABC using alternative algorithms: ABC-SMC

Sequential Monte Carlo (ABC-SMC) is an ABC approach where a sequence of distributions is constructed by gradually decreasing tolerance ϵ . The ABC-SMC algorithm starts by sampling a finite number of parameter sets (particles) from the prior distribution and each intermediate distribution (called a generation) is obtained as a weighted sample from the previous distribution that has been perturbed through a kernel $K(\theta|\theta^*)$.

Pseudocode for an ABC-SMC algorithm are as follows (Toni et al., 2009; McKinley et al., 2018):

1. Set the number of generations G , and the number of particles N .
2. Set the tolerance schedule $\epsilon_1 < \epsilon_2 < \dots < \epsilon_G$. Set the generation indicator $g = 1$.
3. Set particle indicator $i = 1$.
4. If $g = 1$, sample θ^{**} from the prior distribution $P(\theta)$. If $g > 1$, sample θ^* from the previous generation $\{\theta_{g-1}^*\}$ with weights $\{w_{g-1}^*\}$, and perturb the particle to obtain $\theta^{**} \sim K(\theta|\theta^*)$.
5. If $P(\theta^{**}) = 0$, return to step 4.
6. Generate n data sets D_j^{**} from the model using θ^{**} and calculate $\hat{P}(D|D^{**}) = (1/n) \sum_{j=1}^n \mathbb{1}(d(D, D_j^{**}) \leq \epsilon_g)$.
7. If $\hat{P}(D|D^{**}) = 0$, return to step 4.
8. Set $\theta_g^{(i)} = \theta^{**}$ and calculated corresponding weight of the accepted particle i

$$w_g^{(i)} = \begin{cases} \hat{P}(D|D^{**})P(\theta^{**}), & \text{if } g = 1. \\ \frac{\hat{P}(D|D^{**})P(\theta^{**})}{\sum_{j=1}^N w_{g-1}^{(j)} K(\theta_g^{(i)}|\theta_{g-1}^{(j)})}, & \text{if } g > 1. \end{cases}$$

9. If $i < N$, increment $i = i + 1$ and go to step 4.
10. Normalise the weights so that $\sum_{i=1}^N w_g^i = 1$.

11. If $g < G$, set $g = g + 1$, go to step 3.

Here θ^{**} denotes proposed parameter set, and D^{**} is model solution with $\theta = \theta^{**}$.

Factors determining how efficiently the parameter space is explored are the choice of summary statistics, the sequence of tolerance values, the number of generations, the number of simulations for each parameter set, and the choice of perturbation kernel. The algorithm takes into account stochasticity of a model by adjusting weights according to a fraction of model runs which produce a distance smaller or equal to the tolerance value. If the model is deterministic, only one dataset needs to be generated, equivalent to $n = 1$ in step 6.

The behaviour of the algorithm can be assessed by calculating Effective Sample Size (ESS) (Moral et al., 2011):

$$\text{ESS} = \left(\sum_{i=1}^N (w_g^{(i)})^2 \right)^{-1}. \quad (5)$$

The ESS takes values between 1 and N . If ESS value falls below a certain threshold, the algorithm may be deemed to be an inefficient way to sample from the target distribution (Prangle, 2014). In (Moral et al., 2011) this threshold was set as $N/2$.

As mentioned, the above algorithm requires the assignment of a sequence of tolerances. An alternative would be to infer the tolerance by dropping a proportion of the particles with the highest distance values at each generation $g > 1$. The algorithm starts with a tolerance $\epsilon_1 = +\infty$ and the next value of ϵ_t is chosen as the q th quantile of the distribution of accepted metric values, ϵ_{g-1} , at the previous generation (McKinley et al., 2018).

2.3. Choosing perturbation kernel: ABC-SMC MNN

Common choices for a perturbation kernel are the uniform distribution (Toni et al., 2009), or multivariate normal distribution (Kypriaios et al., 2017). A perturbation kernel with a wide variance will stop the algorithm from being stuck in local modes, but will lead to a large number of particles being rejected, and thus cause the algorithm to be inefficient. Hence there is a challenge in finding an optimal perturbation variance that balances efficiency with adequately exploring

the parameter space. In practice adaptive routines are often used such as calculating covariance matrix from the previous generation of accepted particles (Lenormand et al. (2013), Filippi et al. (2013)). In the case of the multivariate normal distribution $K(\theta_g^{(i)}|\theta_{g-1}^{(i)}) = \mathcal{N}(\theta_{g-1}^{(i)}, \Sigma)$, where Σ is the covariance matrix. The covariance matrix has been implemented in a number of ways: as an empirical covariance matrix of all particles from generation $g-1$ (Filippi et al. (2013)), as twice the weighted covariance matrix of the previous generation (Lenormand et al., 2013), or as a covariance matrix calculated using M nearest neighbours (MNN) of the particle $\theta_{g-1}^{(i)}$ (Filippi et al., 2013).

Pseudocode for an ABC-SMC MNN algorithm is as follows (Toni et al., 2009; McKinley et al., 2018):

1. Set the number of generations T , the number of particles N and the number of nearest neighbours M .
2. Set the tolerance schedule $\epsilon_1 < \epsilon_2 < \dots < \epsilon_G$. Set the generation indicator $t = 1$.
3. If $t > 1$, find M nearest neighbours of $\theta_{g-1}^{(i)}$ and calculate empirical covariance matrices $\Sigma(\theta_{g-1}^{(i)}, M)$ for all particles $i = 1 \dots N$.
4. Set particle indicator $i = 1$.
5. If $g = 1$, sample θ^{**} from the prior distribution $P(\theta)$. If $g > 1$, sample θ^* from the previous generation $\{\theta_{g-1}\}$ with weights $\{w_{g-1}\}$, and perturb the particle to obtain $\theta^{**} \sim \mathcal{N}(\theta^*, \Sigma(\theta^*, M))$.
6. If $P(\theta^{**}) = 0$, return to step 4.
7. Generate n data sets D_j^{**} from the model using θ^{**} and calculate $\hat{P}(D|D^{**}) = (1/n) \sum_{j=1}^n \mathbb{1}(d(D, D_j^{**}) \leq \epsilon_g)$.
8. If $\hat{P}(D|D^{**}) = 0$, return to step 4.
9. Set $\theta_g^{(i)} = \theta^{**}$ and calculated corresponding weight of the accepted particle i

$$w_g^{(i)} = \begin{cases} \hat{P}(D|D^{**})P(\theta^{**}), & \text{if } g = 1. \\ \frac{\hat{P}(D|D^{**})P(\theta^{**})}{\sum_{j=1}^N w_{g-1}^{(j)} \mathcal{N}(\theta_g^{(i)}|\theta_{g-1}^{(j)}, \Sigma(\theta_g^{(i)}, M))}, & \text{if } g > 1. \end{cases}$$

10. If $i < N$, increment $i = i + 1$ and go to step 4.
11. Normalise the weights so that $\sum_{i=1}^N w_g^i = 1$.
12. If $g < G$, set $t = t + 1$, go to step 5.

For a case $M = N$, empirical covariance matrix $\Sigma(\theta_{g-1}^{(i)}, M) = \Sigma(\theta_{g-1})$. A normalised Euclidean distance can be used when searching for the nearest neighbours, i.e. differences are divided by the range of the prior of each parameter. This is especially useful if parameter values differ by few orders of magnitude.

3. Case studies

We present the implementation and potential problems that can arise while using ABC for infectious disease models with three case studies. In case study 1 we highlight the parameter identifiable of a deterministic SIR; and in case studies 2 and 3 we investigate the performance of ABC in terms of required number of model runs and effective sample size under few different choices for implementation of algorithm. Code to implement these case studies is written in R and is available to download at https://github.com/amanda-minter/abc_R.

3.1. Case study 1: a simulated epidemic

3.1.1. Data and model

In case study 1, we implement the ABC-rejection algorithm in R. We adapted an existing application of ABC applied to common cold data presented in (Toni et al., 2009). In our application, we simulated an epidemic with known parameters and attempt to re-estimate these parameters using ABC with different resolutions of data. To describe the infection dynamics of the epidemic, we employed a SIR model. In the model, individuals are classed as susceptible (S), infected (and

infectious) (I) or recovered (R).

$$\begin{aligned} \frac{dS}{dt} &= -\beta SI/N \\ \frac{dI}{dt} &= \beta SI/N - \gamma I \\ \frac{dR}{dt} &= \gamma I \end{aligned}$$

where $N = S + I + R$. Daily counts of infected recovered individuals were simulated using the deterministic SIR model with $\beta = 1.5$, $\gamma = 0.5$, giving $R_0 = 3$. The initial conditions used were $S(0) = 99$, $I(0) = 1$ and $R(0) = 0$ and the simulation lasted for 17 days.

We assumed that the data available was (1) collected from the first day of the outbreak (daily counts of infected and recovered individuals) or collected on day 17 only where either the (2) the number of recovered individuals was recorded or (3) the number of infected and recovered individuals was recorded. In all cases, the model simulations were run from day 1 to day 17 based on the length of the daily count data or having the number of infected and/or recovered individuals on day 17.

Note that in this example, if we knew the epidemic had ended, the final size of the epidemic could be calculated using the closed form expression for an SIR model. Here we illustrate the summary statistics that could be calculated using model simulations, assuming that applications in practice would be far more complex than this introductory example.

3.1.2. ABC setup

We implemented three ABC-rejection algorithms in R (Table 1). In the first, we used the daily counts of infected and recovered individuals from day 1 to day 17 to estimate our model parameters. The distance measure was the square root of the sum of squared differences of the number of infected and recovered individuals over time (Table 1). In the second, we assumed that the data consisted of the number of recovered individuals on day 17. Hence, our data is $\mu = R(17)$, and we calculate $\mu^* = R^*(17)$ from our model simulations using the summary statistic i.e. the number of recovered individuals at the end of the model simulation. The distance measure between μ and μ^* was the square root of the squared difference. In the third, we again assumed that the data consisted of the number of recovered individuals on day 17 but in addition, we also calculated the distance between the number of infected individuals on day 17 and model predicted number.

The number of desired particles in all examples was chosen as $N = 1000$. We ran a pilot ABC-rejection algorithm with a smaller N and found that it was difficult to distinguish whether the lack of accuracy was from the small number of particles or the choice of summary statistic. The tolerances were chosen to maximise the precision in estimation of model parameters. For the first example, we ran the ABC-rejection algorithm for a vector of tolerances to choose the tolerance which gave the smallest inter-quartile range of the model parameters (Figure S1). We choose $\epsilon = 20$ as this had the smallest inter-quartile range. In principle, the tolerance could be much lower for the first example, but given the inefficient sampling of the model parameters, the run-time of the algorithm would be long. For examples two and

Table 1

Description of the ABC-rejection algorithms implemented in R. Model simulations are denoted by $*$.

Algorithm	Summary statistic	Distance measure	ϵ
1	–	$\sqrt{\sum_{t=1}^T (I_t - I_t^*)^2} + \sqrt{\sum_{t=1}^T (R_t - R_t^*)^2}$	20
2	$\mu^* = R^*(17)$	$\sqrt{(\mu - \mu^*)^2}$	1
3	$\mu_1^* = R^*(17),$ $\mu_2^* = I^*(17)$	$\sqrt{(\mu_1 - \mu_1^*)^2} + \sqrt{(\mu_2 - \mu_2^*)^2}$	2

Table 2
Definitions and prior distributions of the SIR model parameters.

Parameter	Definition	Prior distribution
β	Transmission rate (day^{-1}).	Uniform (0, 3)
γ	Recovery rate (day^{-1}).	Uniform (0, 1)
S_0	Initial number of susceptibles.	Pois ($\lambda = 100$)

three, we choose the smallest possible tolerance for the model prediction. For example 2, this was $\epsilon = 1$ (comparing the number of recovered individuals on day 17) and example 3, this was $\epsilon = 2$ (comparing the number of recovered individuals and infected individuals on day 17). In all cases, we chose an informative prior for the initial number of susceptibles and uniform priors for the model parameters (Table 2).

3.2. Results

The posterior estimates of the model parameters varied in accuracy and precision for the three ABC-rejection algorithms. When the data consisted of daily counts of infected and recovered individuals (ABC-rejection 1), the true model parameters were retrieved. When the data consisted of the final size of the epidemic (ABC-rejection 2), or of the final size of the epidemic and number of infecteds at the end of the epidemic (ABC-rejection 3) then the posterior distribution of the transmission rate (β) was close to the Uniform prior distribution despite having very low tolerances. The posterior distribution for the recovery rate (γ) had a narrower distribution in ABC-rejection-3. In both ABC-rejection 2 and 3, the posterior of the basic reproduction number $R_0 = \beta/\gamma$ was more accurate than the model parameters.

When the acceptance criteria consisted only of matching the final size of the epidemic (ABC-rejection 2) the accepted model parameters included epidemics which had not ended within the time frame of the true epidemic (Fig. 3). By adding the second criteria of matching the final number of infecteds (0) then we see that the accepted model parameters gave epidemics which had finished within the time frame (Fig. 3). However, the model parameters could still be accepted where the transmission rate (β) was too high or the recovery rate (γ) was too low.

3.3. Case study 2: The 2004 measles outbreak in Malawi

3.3.1. Data and model

In this case study we have compared the performance of the ABC rejection and ABC-SMC algorithms using the 2010 measles outbreak in Malawi data and stochastic model with age structure. This example illustrates challenges to be expected when inferring parameters from a real-world data and a complex epidemic model.

In age-structured models, it is assumed that individuals within a particular age class exhibit more similar behaviour in comparison to other age classes. In the context of epidemiological modelling, infected individual has a higher probability of infecting susceptibles of the same age cohort. This is achieved by defining a Who Acquires Infection From Whom (WAIFW) matrix (Anderson and May, 1982). In this case study, we use the WAIFW matrix for measles from (Shea et al., 2014):

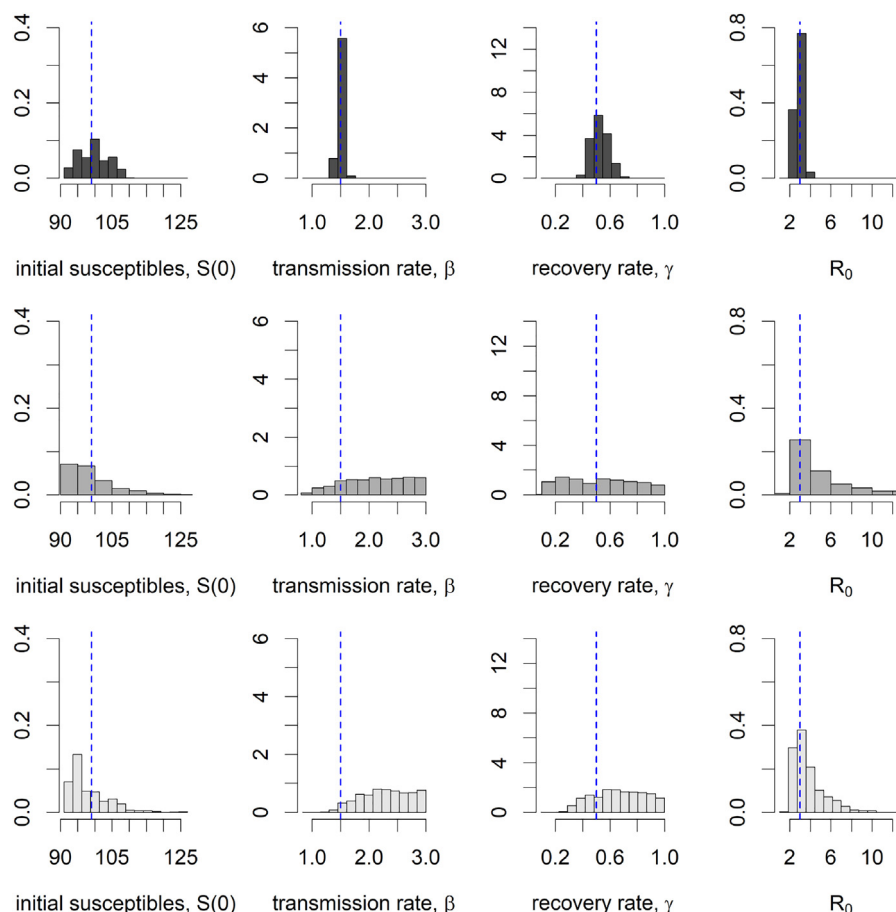


Fig. 2. The posterior distributions of the model parameters for algorithms 1–3 (top row to bottom row). The blue dashed lines indicate the true value of the parameters used to simulate the data. (For interpretation of the references to color in this figure legend, the reader is referred to the web version of this article.)




Fig. 3. The true number of infected (red circle) and recovered individuals (green triangle) and SIR model simulations (shades of grey) using 10 randomly selected parameter sets from the posterior distributions in Fig. 2. (For interpretation of the references to color in this figure legend, the reader is referred to the web version of this article.)

$$A = \begin{bmatrix} 1 & \frac{11}{12} & \frac{5}{6} & \frac{3}{4} & \frac{2}{3} & \dots & \frac{2}{3} \\ \frac{11}{12} & 1 & \frac{11}{12} & \frac{5}{6} & \frac{3}{4} & \dots & \frac{2}{3} \\ \frac{5}{6} & \frac{11}{12} & 1 & \frac{11}{12} & \frac{5}{6} & \dots & \frac{2}{3} \\ \frac{3}{4} & \frac{5}{6} & \frac{11}{12} & 1 & \frac{11}{12} & \dots & \frac{2}{3} \\ \frac{2}{3} & \frac{3}{4} & \frac{5}{6} & \frac{11}{12} & 1 & \dots & \frac{2}{3} \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \frac{11}{12} \\ \frac{2}{3} & \frac{2}{3} & \frac{2}{3} & \frac{2}{3} & \frac{2}{3} & \dots & 1 \end{bmatrix} \quad (6)$$

This is a symmetric matrix with its values decreasing away from the main diagonal. We use a stochastic Susceptible-Exposed-Infected-Recovered (SEIR) model. The force of infection on an individual in age class i at time t is then given by:

$$\lambda_i(t) = \beta \sum_j A_{ij} I_j, \quad (7)$$

where β is transmission rate, I_j is the number of infectious individuals in age class j and A_{ij} is the (i, j) element in the WAIFW matrix.

In this case study we will investigate the 2010 measles outbreak in Malawi (Minetti et al., 2013). The data consist of weekly cases and percentage of cases in three age groups (0.5–5 years, 5–15 and 15+ years). The cumulative number of cases during the outbreak was 134,000. We stratified the population into 22 age classes: 6–12 months, 1–2 years, 2–3 years, ..., 19–20 years, and 20+ years; and assume that the age distribution of susceptible individuals in the population follows the Gamma(a_{sh} , a_{rt}) distribution (Shea et al., 2014). We assume that all individuals 0–6 months old have maternal immunity (Caceres et al., 2000). The model is simulated using the tau-leaping algorithm with time intervals of one day (Gillespie, 2001).

The stochastic compartmental model with age-structure for measles has seven parameters. Duration of latent and infectious periods is seven days (Anderson and May, 1985). Therefore the number of parameters to be estimated is five.

3.3.2. ABC setup

We chose uniform priors for five model parameters (Table 3). As there are two different types of data describing the outbreak (i.e. temporal change in a number of cases and percentage of cases in three age groups), we use an intersection metrics in the ABC-SMC algorithm (McKinley et al., 2018). The distance measure for temporal data is the square root of the sum of squared differences of the number of weekly cases:

$$d^T(D, D^*) = \sqrt{\sum_w (I_w - I_w^*)^2}. \quad (8)$$

Table 3

Definitions and prior distributions of estimated parameters for stochastic compartmental model for measles.

Parameter	Definition	Prior distribution
β	Transmission rate (day^{-1})	Uniform $(0, 5 \times 10^{-6})$
N_0	Susceptible population size before outbreak	Uniform $(140 \times 10^3, 300 \times 10^3)$
f_E	Fraction of initially exposed	Uniform $(0, 0.001)$
a_{sh}	Shape parameter for age distribution	Uniform $(0.8, 1.4)$
a_{rt}	Rate parameter for age distribution	Uniform $(0.1, 0.4)$

The distance measure for age data is the square root of the sum of squared differences of the percentage of infected individuals in the three age groups (0.5–5 years, 5–15 and 15+ years):

$$d^A(D, D^*) = \sqrt{\sum_{a=1}^3 \left(I_a / \left(\sum_{k=1}^3 I_k \right) - I_a^* / \left(\sum_{k=1}^3 I_k^* \right) \right)^2}. \quad (9)$$

We set the pair of thresholds $\{\theta^T, \theta^A\}$ for ABC rejection algorithm in the following way. If we desire the simulated number of cases to be within $\pm 20\%$ of the observed number of cases during the outbreak duration (i.e. 54 weeks), we get $d^T(D, D^*) = \sqrt{\sum_{w=1}^{54} (0.4 \times I_w^*)^2} = 11,900$. In a similar way, if we require the differences in each age group to be $\pm 1\%$, this gives $d^A(D, D^*) = \sqrt{3 \times 2^2} = 2.82$. Therefore we set $\epsilon^T = 10,000$ and $\epsilon^A = 3$ for the ABC-rejection algorithm.

Next, we will investigate the performance of ABC-SMC algorithm. We will explore two strategies: (i) calculating covariance matrix using all sampled particles from the previous generation (ABC-SMC); and (ii) calculating covariance matrix using M nearest neighbours of the particle $\theta_{g-1}^{(i)}$ (ABC-SMC MNN) (Filippi et al., 2013).

We set the same schedule of tolerances to run both the ABC-SMC and ABC-SMC MNN algorithms. For the first generation, we set the pair of thresholds in the following way: as $d^T(D, D^*) = \sqrt{\sum_{w=1}^{54} (2 \times I_w^*)^2} = 59,500$ (i.e. we allow the discrepancy to be twice as large as weekly number of cases) and $d^A(D, D^*) = \sqrt{3 \times 10^2} = 17.32$ (i.e. we allow the discrepancy to be as large as ten percent for each age group), we set $\epsilon_1^T = 50,000$ and $\epsilon_1^A = 15$. Then we gradually decreased the threshold levels until they reached the same fixed level as was used for the ABC rejection method.

The number of desired particles was set to $N = 1000$ and the number of generation to $G = 10$. The tolerance values for the last generation was set to be the same as for the ABC rejection algorithm. We have investigated how ABC-SMC and ABC-SMC MNN algorithms performed when a number of model runs for each parameter set was $n = 1$, $n = 10$ and $n = 100$. For ABC-SMC MNN method we set the number of neighbours to $M = 50$ and $M = 100$.

Finally, we tested the performance of ABC-SMC and ABC-SMC MNN by setting the new tolerances ϵ_g as a median (i.e. 50th quantile) of

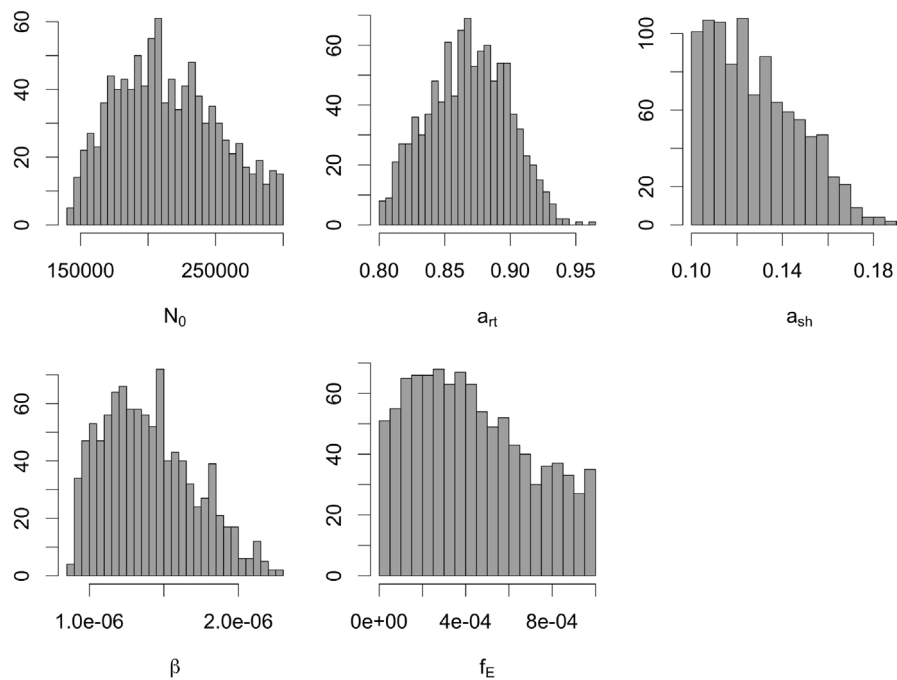


Fig. 4. The posterior distributions of parameters obtained using the ABC rejection algorithm.

$d^m(D, D^*) < \epsilon_{g-1}^m$ (McKinley et al., 2018). Again, the final generation tolerance level was set to be equal or less than values used for the ABC rejection method.

3.3.3. Results

First, we have employed the ABC rejection algorithm with tolerances $\epsilon^T = 10,000$ and $\epsilon^A = 3$. The number of model runs required to obtain 1000 accepted particles was 29, 224, 520. This gives that the acceptance rate for the ABC rejection approach is 0.00003. The posterior estimates for the model parameters are shown in Fig. 4(A)–(E). We have estimated mean value for the transmission rate equal to 1.39×10^{-6} per day and range as $(8.7 \times 10^{-7}, 2.26 \times 10^{-6})$.

Model simulations using 10 randomly selected parameter sets from the posterior distribution are shown in Fig. 5. Demographic distribution of cases was calculated using these 10 simulations. There is a good agreement between the observed data and simulated data for both temporal dynamics of the outbreak and fraction of cases in three age groups.

Applying the ABC-SMC algorithm for $G = 10$ generations required one hundred times fewer model runs than the ABC-rejection algorithm: the total number of model runs required to obtain 1000 accepted

particles for the ABC-SMC was 267, 388. We have repeated parameter inference by re-running model 10 and 100 times for each proposed parameter set. The total number model runs to obtain 1000 accepted values for each generation is given in Table 4. For the first generation, where parameters are sampled from the prior distribution, the acceptance rate was higher for multiple runs of the model (33.8% for $n = 10$ and 39.5% for $n = 100$) in comparison to a single model run (28.6%). The acceptance rates for generations 2–10 were as follows: $n = 1$ from 1.7% to 8.1% (ABC-SMC) and from 2.9% to 14.7% (ABC-SMC MNN); for $n = 10$ from 2.9% to 9.8% (ABC-SMC) and from 4.8% to 16.5% (ABC-SMC MNN); for $n = 100$ from 3.0% to 10.3% (ABC-SMC) and from 6.3% to 18.8% (ABC-SMC MNN). This result suggests that increasing a number of runs for each proposed parameter set increased parameter acceptance rate.

Using the multivariate normal kernel with 50 nearest neighbours increased the performance of the ABC-SMC algorithm. For example, the scheme required 160, 644 model runs in total in order to accept 1000 particles over 10 generations for $n = 1$, which gives around 40% reduction in a number of simulations. When a number of nearest neighbours was set to 100, the ABC-SMC MNN algorithm showed improvement with respect to total number of runs in comparison to ABC-SMC,




Fig. 5. Comparing observed outbreak data (Minetti et al., 2013) with model simulations: number of cases each week (left) and percentage of cases in three age groups (right).

Table 4Number of model runs needed to accept $N = 1000$ particles in each generation, acceptance rate and ESS for ABC algorithms run with a fixed set of tolerances.

Generation	ϵ_g^T	ϵ_g^A	ABC SMC ($M = 1000$)			ABC-SMC MNN					
			$n = 1$	$n = 10$	$n = 100$	$M = 100$			$M = 50$		
						$n = 1$	$n = 10$	$n = 100$	$n = 1$	$n = 10$	$n = 100$
1	50,000	15	3493	27,120	271,900	3616	27,660	267,300	3465	27,880	131,600
2	27,500	10	12,331	101,440	966,100	9819	79,680	657,400	8886	62,550	622,200
3	25,000	9	13,710	117,290	983,300	9819	77,890	756,500	6779	60,580	530,900
4	22,500	8	16,917	122,300	1,129,600	11,535	93,350	930,600	8242	75,230	533,200
5	20,000	7	18,597	144,820	1,236,300	15,593	114,420	852,100	10,396	86,870	604,000
6	17,500	6	23,592	193,140	1,612,600	18,819	130,380	1,257,200	13,120	94,650	733,100
7	15,000	5	33,657	234,250	2,178,200	24,741	189,250	1,448,800	17,360	143,730	1,219,300
8	12,500	4	58,158	339,350	3,269,300	52,621	224,640	2,241,500	27,582	207,310	1,565,600
9	11,000	3.5	43,806	301,010	2,287,500	42,915	220,460	1,904,500	31,138	187,620	1,442,500
10	10,000	3	43,127	286,220	2,135,700	65,852	278,600	2,227,900	33,676	199,510	1,182,700
Total			267,388	1,869,420	16,070,500	255,372	1,436,330	12,543,800	160,644	1,145,830	8,565,100
Acceptance rate			0.037	0.053	0.062	0.039	0.069	0.079	0.062	0.087	0.116
ESS			853	606	444	705	587	476	687	623	366

but the performance was much lower when $M = 50$.

We have investigated how ESS for the last generation compares for different configurations of the ABC algorithm. We found that the ESS values decreased when (i) number of nearest neighbours decreased; and (ii) number of model runs for each proposed parameter set increased. For the runs with $n = 100$, the ESS was below 500. In contrast, acceptance rate calculated over all generations has increased when increasing number of model runs. Therefore, there is a trade-off between number of model runs, acceptance rate and variance of the weights of accepted particles when choosing a number of runs per proposed parameter set and number of nearest neighbours for covariance matrix of the perturbation kernel. Our analysis based on both the ABC-SMC and ABC-SMC MNN indicates that having $n > 1$ increases computation burden and decreases the ESS, which cannot be out-weighted by improvement in the overall acceptance rate. This result is in agreement with McKinley et al. (McKinley et al., 2009), who found that for a given tolerance level, increasing the number of repeats did not produce markedly more precise posteriors.

Finally, we have explored the performance of the ABC-SMC and ABC-SMC MNN for the case when the tolerance is calculated as the median from the previous generation. Total number of model runs needed to accept $N = 1000$ particles in each generation is given in Table 5. All approaches required $G = 7$ generations to reach $\epsilon_g^T < 10,000$ and $\epsilon_g^A < 3.0$. Using 50 nearest neighbours when calculating covariance matrix for perturbation kernel required 40% less model runs. Again, the increase in acceptance rate came at a cost of decreasing ESS.

Table 5Number of model runs needed to accept $N = 1000$ particles in each generation, acceptance rate and ESS for ABC algorithms run with tolerances calculated from previous generation ($n = 1$).

Generation	ABC-SMC			ABC-SMC MNN					
	ϵ_g^T	ϵ_g^A		$M = 100$			$M = 50$		
				ϵ_g^T	ϵ_g^A		ϵ_g^T	ϵ_g^A	
2	67,160	8.21	4,404	62,885	8.45	4705	67,160	8.21	4450
3	29,749	5.39	11,301	29,749	5.38	7696	30,705	5.34	6405
4	27,490	4.20	29,494	26,422	4.11	18,456	26,612	4.07	16,843
5	17,837	3.59	40,658	17,585	3.34	25,863	16,965	3.40	20,476
6	12,550	3.14	74,170	12,514	2.90	52,082	12,199	2.93	35,891
7	9687	2.84	84,400	9591	2.68	71,236	9782	2.67	61,315
Total			245,427			181,038			146,380
Acceptance rate			0.028			0.038			0.047
ESS			861			750			682

3.4. Case study 3: Citrus tristeza virus spread in an orchard

3.4.1. Data and model

In this case study we will investigate the data based on small scale field observations of citrus tristeza virus (CTV) spread in an orchard (Marcus et al., 1984). The trees were arranged in a rectangular frame, with an inter-row distance of 5.5 m and a between-column distance of 4 m. There were 131 infected trees in 1981 and 45 infected trees in 1982. Maps of susceptible and infected trees in 1981 and 1982 are shown in Fig. 6. The virus affects citrus trees and is transmitted by the brown citrus aphid (Keeling and Rohani, 2007).

We model the spread of CTV using a spatio-temporal individual-based stochastic model. Each tree is classified as susceptible or infected. The force of infection on a susceptible tree i at time t is given by (Marcus et al., 1984; Keeling et al., 2004):

$$\lambda_i(t) = \frac{\beta}{365} \sum_{j \in \text{infected}} K(d_{ij}), \quad (10)$$

where β is yearly transmission rate, d_{ij} is the distance between trees i and j , and $K(d) = d^{-2\alpha}$ is a distance kernel. We have simulated a number of new cases every day for a year starting with the trees infected in 1981.

3.4.2. ABC setup

We implemented ABC algorithm with the distance measure which compares observed and simulated distribution of a minimal distance between trees infected in 1981 and 1982:




Fig. 6. Locations of trees: susceptible trees (green circles) and trees infected by CTV (red circles) (Marcus et al., 1984). (For interpretation of the references to color in this figure legend, the reader is referred to the web version of this article.)

$$d^S(D, D^*) = \sqrt{\sum_{\delta \in \Delta} (n_{\text{obs}}(\delta) - n_{\text{sim}}(\delta))^2}, \quad (11)$$

where Δ is a set of all distances between trees. There are 992 unique distances between all trees. A number of observed and simulated minimal distances is calculated as:

$$n(\delta) = \sum \mathbb{1}(\min \|a - b\| = \delta), \quad (12)$$

where a are the coordinates of trees infected in 1981 (i.e. observed incidences of disease); and b are the coordinates of trees infected in 1982 (i.e. observed or simulated incidences of disease).

The number of desired particles was $N = 1000$. If we assume that there were no new cases in 1982, we get that $d^S(D, D^*) = 20$. Therefore we chose $\epsilon_1 = 20$. There were 45 infected trees in 1982, therefore we required that the threshold for the last generation would be less than $\sqrt{45/2} = 4.74$. The schedule of tolerances was then set to $\epsilon = 20, 15, 10, 7.5, 6, 5, 4.5$.

We have investigated how ABC SMC MNN algorithm performed when covariance matrix of the perturbation kernel was calculated using $M = 1000$, $M = 500$, $M = 100$ and $M = 50$ of nearest neighbours. Parameter priors are given in Table 6.

3.4.3. Results

First, we implemented ABC-SMC MNN with $M = 50$. Fig. 7 shows the posterior distributions of the model parameters and fitted distance kernel. We estimated the mean of power law decay parameter to be 1.32. The value previously identified by MCMC analysis is $\alpha = 1.3$ (Gibson, 1997) and using conservation of pattern method $\alpha = 1.35$ (Keeling et al., 2004).

Model simulations using 100 randomly selected parameter sets from the posterior distribution are shown in Fig. 8. It can be seen that the mean number of simulated infections within a year (41.5 cases) is close to the observed number of new infections (45 cases). The distributions of simulated number of minimal distances lie close to the observed number of minimal distances, as expected. Figure on the bottom shows a proportion of simulations which had a tree being infected at the start of 1982 with darker shades illustrating a higher number of simulations.

Performance of ABC SMC MNN algorithm implemented using $M = 1000$, $M = 500$, $M = 100$ and $M = 50$ is given in Table 7. The mean of power law decay parameter decreased from 1.39 to 1.35 when M decreased. We have found that decreasing number of neighbour particles decreases number of model runs required to accept 1000 particles and increases parameter acceptance rate. However, the ESS

values for the last generation were at the similar level (630–650).

For a comparison, we have employed the ABC rejection algorithm with tolerance $\epsilon = 4.5$. The mean of power law decay parameter was 1.39, i.e. similar to the result with $M = 1000$. The number of model runs required to obtain 1000 accepted particles was 3, 820, 136. This gives that the acceptance rate for the ABC rejection approach is 0.00026.

4. Discussion

The purpose of this paper was to provide an introductory tutorial in using ABC for infectious disease modelling. We have introduced the user choices required to implement the ABC-rejection algorithm and the ABC-SMC algorithm. These user defined choices can be model or data based. In addition, with larger data sets and more complex models, there are also computational trade-offs to be made, such as with a more sophisticated ABC algorithm. Here, we have presented three case studies to illustrate how to implement ABC-rejection and ABC-SMC in R. However, there are numerous applications and examples which we have not covered but have been by other authors. We have applied the ABC-SMC algorithm to stochastic models, but the technique has been used for a system of ordinary differential equations as well (Toni et al., 2009; Filippi et al., 2013).

In case study 1 we illustrated the information lost and the changes that need to be made to the algorithm based on different resolutions of data. Daily counts of infected and recovered individuals provides enough information to re-estimate model parameters. However, when the number of recovered or infected individuals was known, only the basic reproduction number can be re-estimated and with less precision. This is due to the lack of temporal information in the data and the correlation between the transmission rate and recovery rate. If the user was able to specify an informative prior for either the transmission rate or recovery rate, then the other parameter could be estimated.

The ABC-rejection algorithm has inefficient sampling in comparison to the ABC-SMC algorithm. By re-sampling from a smaller generation and slowly decreasing tolerance, we saw that in case study 2, ABC-SMC required hundred times fewer model runs than the ABC-rejection algorithm: the number of model runs required to obtain 1000 accepted particles for ABC rejection algorithm was 29, 224, 520, while for ABC-SMC it was 267, 388. Furthermore, if only 50 nearest neighbours of particles were used to parametrise the perturbation kernel, the number of model runs was reduced to 160, 644. Running the model for each proposed parameter set more times resulted in a higher acceptance rate. However, the trade-off between running model multiple times and increasing acceptance rate will depend on a cost of each model run and decreased effective sample size. For case study 3, decreasing number of neighbour particles decreased number of model runs and increased parameter acceptance rate, but had no effect on values of effective sample size.

The number of desired particles in all case studies was set to

Table 6

Definitions and prior distributions of estimated parameters.

Parameter	Definition	Prior distribution
α	Power law decay.	Uniform (0, 5)
β	Transmission rate (year^{-1})	Uniform (0, 10)

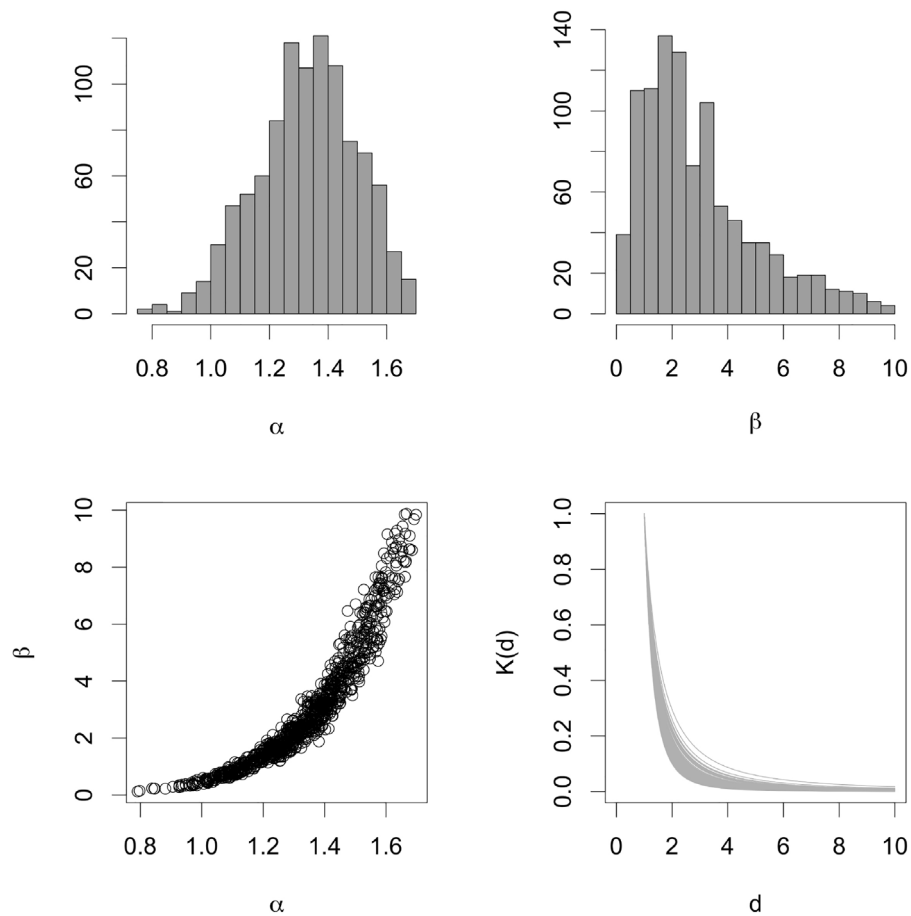


Fig. 7. The posterior distributions of the model parameters and fitted distance kernel.

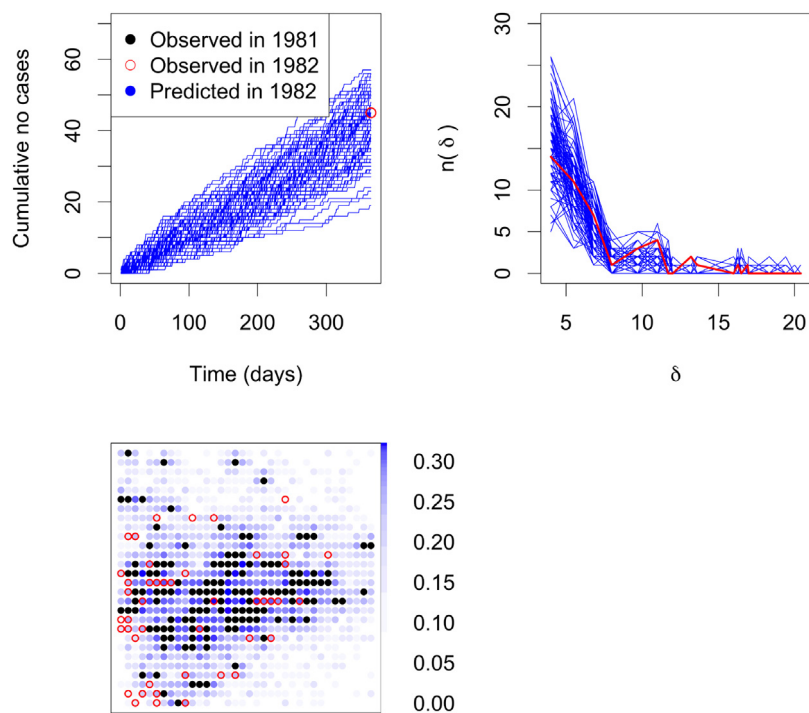


Fig. 8. Comparing observed outbreak data (Marcus et al., 1984) with model simulations: number of cumulative cases each day (top left), number of observed and simulated minimal distances (top right), and spatial distribution (bottom left).

Table 7
Comparison of ABC SMC MNN algorithms.

	$M = 1000$	$M = 500$	$M = 100$	$M = 50$
$E[\lambda]$	1.39	1.35	1.34	1.32
$E[\beta]$	3.83	3.11	3.10	2.95
No runs	155,108	108,590	95,116	89,863
Accept. rate	0.045	0.064	0.073	0.077
ESS	650	633	644	630

$N = 1000$ as in (Toni et al., 2009). The choice of number of particles will affect the speed of the ABC algorithm and the accuracy of the final posterior distribution. Additionally, the magnitude of the tolerance will also affect the speed of the algorithm and accuracy of the algorithm. The choice of tolerance in the ABC-rejection algorithm can be informed by trial and error. In the ABC-SMC algorithm, instead of proposing a vector of tolerances, the tolerance at each population can be calculated as the α quantile of the distances of the accepted particles (Prangle, 2017). The user then has to choose some quantile (q) and the number of populations (G). This quantile method means that the user does not need to supply a vector of tolerances, however must choose the number of populations (G). Convergence can be assessed by the difference using the inter-quartile ranges of the values of accepted particles as a measure of goodness of fit between successive intermediate distributions (Toni et al., 2009).

Exploration of the data is also essential to inform the choice of summary statistics. In all cases we have used a distance measure as the squared difference between observed and simulated data. Other options are available as the absolute difference (Sunnåker et al., 2013). If the model output is a two peak epidemic then an appropriate summary statistic may be to match the time of the two peaks, or the ratio of the size of one peak to another. Distance can include summary statistics which is a function of model outputs only, for example, simulated infection pressures (Prangle et al., 2018). Proposed solutions to choosing optimal summary statistics for ABC include a sequential scheme to assess the information content of a list of summary statistics (Joyce and Marjoram, 2008), minimisation of the estimated entropy of the posterior approximation (Nunes and Balding, 2010), finding statistical summaries of the data that minimise the loss of information (Barnes et al., 2012), constructing summary statistics as estimates of the posterior mean of the parameters (Fearnhead and Prangle, 2012), or employing a machine learning tool named random forests where a reference table generated from a prior distribution is used to train regression trees (Pudlo et al., 2015; Raynal et al., 2018). We recommend looking at the literature for inspiration the user defined choices including the summary statistics. The type of data available and problems which arise from infectious disease modelling reoccur in the literature. For a review of applications of ABC see Beaumont (2010) and the papers highlighted in the discussion. Also for applications of ABC-SMC in infectious disease modelling see Conlan et al. (2012), Hladish et al. (2018).

In this paper, we have assumed that the user is writing their ABC algorithm in computer code. There are R packages which implement ABC algorithms: **abc** (ABC-rejection method) (Csilléry et al., 2012), **abctools** (ABC-rejection method) (Nunes and Prangle, 2015), **abcrf** (ABC parameter inference via random forests) (Raynal et al., 2018), and **easyABC** (ABC-rejection and ABC-SMC methods) (Jabot et al., 2013). While learning the steps of ABC, the authors believe that the user is more likely to understand each step of the algorithm with code, in comparison to a package. In addition, by writing the full ABC algorithm in code, the user has the flexibility to alter steps according to their type of model and their data. However, there are some simple applications of ABC where a package would be appropriate if that is the user's preference. We hope that the case studies we presented here can give practically useful insight when making various ABC algorithm tuning choices.

ABC is a powerful tool but an ABC algorithm must be tailored to the model and data at hand. The performance of any model fitting algorithm is restricted by the available data which the user is trying to fit their model to. The user must specify a model appropriate for their data source, i.e. parameters must all be identifiable from the data at hand. If the model is stochastic, investigating the variability in the model output for the same set of parameter values will inform how the choice of n , the number of times to run the model and also what summary statistics to choose. By simulating data from the model, and attempting to re-estimate the known parameters the user can identify changes that are needed for their question specific ABC algorithm.

Conflicts of interest

The authors have no conflicts of interest to declare.

Acknowledgements

We would like to thank Will Probert for his comments on the manuscript and the reproducibility of the code. The authors thank the UK National Institute for Health Research Health Protection Research Unit (NIHR HPRU) in Modelling Methodology at Imperial College London in partnership with Public Health England (PHE) for funding (grant HPRU-2012–10080).

Appendix A. Supplementary data

Supplementary data associated with this article can be found, in the online version, at <https://doi.org/10.1016/j.epidem.2019.100368>.

References

- Anderson, R., May, R., 1982. Control of communicable disease by age-specific immunisation schedules. *Lancet* 319 (8264), 160.
- Anderson, R.M., May, R.M., 1985. Age-related changes in the rate of disease transmission: implications for the design of vaccination programmes. *J. Hyg.* 94 (03), 365–436.
- Anderson, R.M., May, R.M., 1991. *Infectious Diseases of Humans*, 1st ed. Oxford University Press.
- Babad, H.R., Nokes, D.J., Gay, N.J., Miller, E., Morgan-Capner, P., Anderson, R.M., 1995. Predicting the impact of measles vaccination in England and Wales: model validation and analysis of policy options. *Epidemiol. Infect.* 114 (02), 319.
- Baragatti, M., Pudlo, P., 2014. An overview on approximate Bayesian computation. *ESAIM: PROCEEDINGS* 44, 291–299.
- Barnes, C.P., Filippi, S., Stumpf, M.P.H., Thorne, T., 2012. Considerate approaches to constructing summary statistics for ABC model selection. *Stat. Comput.* 22 (6), 1181–1197.
- Beaumont, M.A., 2010. Approximate Bayesian computation in evolution and ecology. *Annu. Rev. Ecol. Evol. Syst.* 41 (1), 379–406.
- Beaumont, M.A., Zhang, W., Balding, D.J., 2002. Approximate Bayesian computation in population genetics. *Genetics* 162 (4), 2025–2035.
- Caceres, V.M., Strebel, P.M., Sutter, R.W., 2000. Factors determining prevalence of maternal antibody to measles virus throughout infancy: a review. *Clin. Infect. Dis.* 31 (1), 110–119.
- Campbell, P.T., McCaw, J.M., McVernon, J., 2015. Pertussis models to inform vaccine policy. *Hum. Vaccines Immunother.* 11 (3), 669–678.
- Chapman, L.A.C., Jewell, C.P., Spencer, S.E.F., Pellis, L., Datta, S., Chowdhury, R., Bern, C., Medley, G.F., Hollingsworth, T.D., 2018. The role of case proximity in transmission of visceral leishmaniasis in a highly endemic village in Bangladesh. *PLOS Negl. Trop. Dis.* 12 (10), e0006453.
- Conlan, A.J.K., McKinley, T.J., Karolemeas, K., Pollock, E.B., Goodchild, A.V., Mitchell, A.P., Birch, C.P.D., Clifton-Hadley, R.S., Wood, J.L.N., 2012. Estimating the hidden burden of bovine tuberculosis in Great Britain. *PLOS Comput. Biol.* 8 (10), 1–14.
- Csilléry, K., François, O., Blum, M.G.B., 2012. abc: an r package for approximate Bayesian computation (ABC). *Methods Ecol. Evol.* 3 (3), 475–479.
- Fearnhead, P., Prangle, D., 2012. Constructing summary statistics for approximate Bayesian computation: semi-automatic approximate Bayesian computation. *J. R. Stat. Soc. Ser. B (Stat. Methodol.)* 74 (3), 419–474.
- Filippi, S., Barnes, C.P., Cornebise, J., Stumpf, M.P., 2013. On optimality of kernels for approximate Bayesian computation using sequential Monte Carlo. *Stat. Appl. Genet. Mol. Biol.* 12 (1).
- Gelman, A., Carlin, J.B., Stern, H.S., Rubin, D.B., 2013. *Bayesian Data Analysis*, 3rd ed. Chapman and Hall/CRC.
- Gibson, G.J., 1997. Markov chain monte carlo methods for fitting spatiotemporal stochastic models in plant epidemiology. *J. R. Stat. Soc. Ser. C (Appl. Stat.)* 46 (2), 215–233.
- Gillespie, D.T., 2001. Approximate accelerated stochastic simulation of chemically

- reacting systems. *J. Chem. Phys.* 115 (4), 1716–1733.
- Hartig, F., Calabrese, J.M., Reineking, B., Wiegand, T., Huth, A., 2002. Statistical inference for stochastic simulation models – theory and application. *Ecol. Lett.* 14 (8), 816–827.
- Hill, E.M., House, T., Dhingra, M.S., Kalpravidh, W., Morzaria, S., Osmani, M.G., Yamage, M., Xiao, X., Gilbert, M., Tildesley, M.J., 2017. Modelling h5n1 in bangladesh across spatial scales: model complexity and zoonotic transmission risk. *Epidemics* 20, 37–55.
- Hladish, T.J., Pearson, C.A.B., Patricia Rojas, D., Gomez-Dantes, H., Halloran, M.E., Vazquez-Prokopec, G.M., Longini, I.M., 2018. Forecasting the effectiveness of indoor residual spraying for reducing dengue burden. *PLOS Negl. Trop. Dis.* 12 (6), 1–16.
- Jabot, F., Faure, T., Dumoulin, N., 2013. EasyABC: performing efficient approximate Bayesian computation sampling schemes using r. *Methods Ecol. Evolut.* 4 (7), 684–687.
- Joyce, P., Marjoram, P., 2008. Approximately sufficient statistics and Bayesian computation. *Stat. Appl. Genet. Mol. Biol.* 7 (1).
- Kanaan, M.N., Farrington, C.P., 2005. Matrix models for childhood infections: a Bayesian approach with applications to rubella and mumps. *Epidemiol. Infect.* 133 (06), 1009.
- Keeling, M.J., Brooks, S.P., Gilligan, C.A., 2004. Using conservation of pattern to estimate spatial parameters from a single snapshot. *Proc. Natl. Acad. Sci. USA* 101 (24), 9155–9160.
- Keeling, M.J., Rohani, P., 2007. *Modeling Infectious Diseases in Humans and Animals*, 1st ed. Princeton University Press, Princeton, N.J.
- Keeling, M.J., Woolhouse, M.E.J., Shaw, D.J., Matthews, L., Chase-Topping, M., Haydon, D.T., Cornell, S.J., Kappey, J., Wilesmith, J., Grenfell, B.T., 2001. Dynamics of the 2001 UK foot and mouth epidemic: stochastic dispersal in a heterogeneous landscape. *Science* 294 (5543), 813–817.
- Kypraios, T., Neal, P., Prangle, D., 2017. A tutorial introduction to Bayesian inference for stochastic epidemic models using approximate Bayesian Computation. *Math. Biosci.* 287, 42–53.
- Lenormand, M., Jabot, F., Deffuant, G., 2013. Adaptive approximate Bayesian computation for complex models. *Comput. Stat.* 28 (6), 2777–2796.
- Marcus, R., Fishman, S., Talpaz, H., Salomon, R., Bar-Joseph, M., 1984. On the spatial distribution of citrus tristeza virus disease. *Phytoparasitica* 12 (1), 45–52.
- Marjoram, P., Molitor, J., Plagnol, V., Tavaré, S., 2003. Markov chain Monte Carlo without likelihoods. *Proc. Natl. Acad. Sci. USA* 100 (26), 15324–15328.
- McKinley, T., Cook, A.R., Deardon, R., 2009. Inference in epidemic models without likelihoods. *Int. J. Biostat.* 5 (1).
- McKinley, T.J., Vernon, I., Andrianakis, I., McCreesh, N., Oakley, J.E., Nsubuga, R.N., Goldstein, M., White, R.G., 2018. Approximate Bayesian computation and simulation-based inference for complex stochastic epidemic models. *Stat. Sci.* 33 (1), 4–18.
- Minetti, A., Kagoli, M., Katsulukuta, A., Huerga, H., Featherstone, A., Chiotcha, H., Noel, D., Bopp, C., Sury, L., Fricke, R., Iscla, M., Hurtado, N., Ducomble, T., Nicholas, S., Kabuluzi, S., Grais, R.F., Luquero, F.J., 2013. Lessons and challenges for measles control from unexpected large outbreak, Malawi. *Emerg. Infect. Dis.* 19 (2), 202–209.
- Moral, P.D., Doucet, A., Jasra, A., 2011. An adaptive sequential monte carlo method for approximate Bayesian computation. *Stat. Comput.* 22 (5), 1009–1020.
- Nunes, M.A., Balding, D.J., 2010. On optimal selection of summary statistics for approximate Bayesian computation. *Stat. Appl. Genet. Mol. Biol.* 9 (1).
- Nunes, M.A., Prangle, D., 2015. abctools: an R package for tuning approximate Bayesian computation analyses. *R Journal* 7 (2), 189–205.
- Perez-Lezaun, A., Pritchard, J.K., Seielstad, M.T., Feldman, M.W., 1999. Population growth of human Y chromosomes: a study of Y chromosome microsatellites. *Mol. Biol. Evolut.* 16 (12), 1791–1798.
- Prangle, D., 2014. Lazy ABC. *Stat. Comput.* 26 (1–2), 171–185.
- Prangle, D., 2017. Adapting the abc distance function. *Bayesian Anal.* 12 (1), 289–309.
- Prangle, D., Everitt, R.G., Kypraios, T., 2018. A rare event approach to high-dimensional approximate Bayesian computation. *Stat. Comput.* 28 (4), 819–834.
- Pudlo, P., Marin, J.-M., Estoup, A., Cornuet, J.-M., Gautier, M., Robert, C.P., 2015. Reliable ABC model choice via random forests. *Bioinformatics* 32 (6), 859–866.
- Raynal, L., Marin, J.-M., Pudlo, P., Ribatet, M., Robert, C.P., Estoup, A., 2018. ABC random forests for Bayesian parameter inference. *Bioinformatics*.
- Sadegh, M., Vrugt, J.A., 2014. Approximate Bayesian computation using markov chain monte carlo simulation: dream(abc). *Water Resour. Res.* 50 (8), 6767–6787.
- Schuette, M., 1999. Modeling the effects of varicella vaccination programs on the incidence of chickenpox and shingles. *Bull. Math. Biol.* 61 (6), 1031–1064.
- Shea, K., Tildesley, M.J., Runge, M.C., Fonnesbeck, C.J., Ferrari, M.J., 2014. Adaptive management and the value of information: learning via intervention in epidemiology. *PLOS Biol.* 12 (10), e1001970.
- Sunnåker, M., Busetto, A.G., Numminen, E., Corander, J., Foll, M., Dessimoz, C., 2013. Approximate Bayesian computation. *PLOS Comput. Biol.* 9 (1), 1–10.
- Toni, T., Welch, D., Strelkowa, N., Ipsen, A., Stumpf, M.P., 2009. Approximate Bayesian computation scheme for parameter inference and model selection in dynamical systems. *J. R. Soc. Interface* 6 (31), 187–202.
- Visser, M.D., McMahon, S.M., Merow, C., Dixon, P.M., Record, S., Jongejans, E., 2015. Speeding up ecological and evolutionary computations in R; essentials of high performance computing for biologists. *PLOS Comput. Biol.* 11 (3), 1–11.
- Wilson, G., Aruliah, D.A., Brown, C.T., Chue Hong, N.P., Davis, M., Guy, R.T., Haddock, S.H.D., Huff, K.D., Mitchell, I.M., Plumbly, M.D., Waugh, B., White, E.P., Wilson, P., 2014. Best practices for scientific computing. *PLOS Biol.* 12 (1), 1–7.